

Figure 1: Even when GT is in the top-5 by embedding, LLM selection drops below a no-LLM argmax baseline under hard distractors
(TB-validated, 62 tasks; claude-sonnet-4-5; baselines: argmax=53.2%, random-from-top-5=31.3%)

