

Skill-selection eval - results log

2026-05-03: SkillsBench full grid

Run: testsets/runs/2026-05-03-grid-full.jsonl Config: 89 tasks x {5, 50, 500} pool size x {random, hard, easy, none} noise x seed 0 = **1068 trials, 0 errors**. Backend: claude -p Max plan, model claude-sonnet-4-5, conc=6, ~48 min wall-clock. Pool builder: max_gt=1 (subsample to 1 GT/task), GT-alias-aware ClawHub exclusion, post-shuffle GT positions recorded.

size	noise	n	Hit@1	Recall	MRR	refusal
5	random	89	80.90%	80.90%	0.809	17.98%
5	hard	89	59.55%	60.67%	0.601	10.11%
5	easy	89	85.39%	85.39%	0.854	14.61%
5	none	89	-	-	-	**97.75%**
50	random	89	83.15%	83.15%	0.831	15.73%
50	hard	89	55.06%	55.06%	0.551	12.36%
50	easy	89	80.90%	80.90%	0.809	17.98%
50	none	89	-	-	-	**95.51%**
500	random	89	76.40%	76.40%	0.764	20.22%
500	hard	89	**41.57%**	41.57%	0.416	12.36%
500	easy	89	71.91%	71.91%	0.719	28.09%
500	none	89	-	-	-	82.02%

Headline findings

1. **Selection collapse under hard noise.** Hard-neg distractors degrade Hit@1 from 59.6% (sz=5) -> 41.6% (sz=500), an 18 pp drop. The model genuinely cannot distinguish GT from its tightest embedding neighbours at scale. 2. **Random / easy noise are mostly survivable.** ~10 pp drop from sz=5 to sz=500 - the model handles uncorrelated distraction. 3. **Refusal in no-GT condition is robust at small/medium sizes (>=95%) but breaks at sz=500 (82%).** The model "hallucinates" a pick in 18% of large no-GT pools - a clear retrieval-system safety concern. 4. **Easy noise at sz=5 gives the highest Hit@1 (85.4%)** - when distractors are blatantly wrong, GT pops out cleanly. This is the "ceiling" reference for the harness.

Position-bias sanity (separate run, sz=500 random, n=9/bin)

Hit@1: early 100% / middle 78% / late 89%. Mild middle-position dip; no catastrophic attention collapse - confirms the 500-pool numbers are not a positional artifact.

Caveats

- * max_gt=1 simplifies but ignores multi-GT tasks (some SkillsBench tasks legitimately need 3-5 skills).
- * Single seed; variance not estimated. Re-run with seeds {0,1,2} for confidence intervals before paper.
- * 25.5% of GT skills have a ClawHub alias at sim >= 0.9; the pool builder excludes the aliased corpus ID from noise, so misses in hard mode at the alias boundary represent the model flipping between embedding-equivalent skills (not a bug - that is the failure mode being measured).

2026-05-03: Terminal-Bench pilot (P4 scope: no execution, no validated GT)

Run: testsets/runs/2026-05-03-tb-pilot.jsonl Config: 30 TB tasks x {5, 50, 500} pool size x {topk_only, random, hard, easy, none} noise x seed 0 = **450 trials, 0 errors**. Pool: top-5 FAISS retrievals from ClawHub serve as the proxy candidate set; treated as "intended" rather than validated GT (P4 dropped harbor execution).

size	noise	n	pick	in_topk	top1=ret	refusal
------	-------	---	------	---------	----------	---------

5	topk_only	30	33.3%	33.3%	16.7%	66.7%
5	random	30	33.3%	33.3%	20.0%	66.7%
5	hard	30	26.7%	26.7%	20.0%	73.3%
5	easy	30	33.3%	33.3%	20.0%	66.7%
5	none	30	0.0%	-	-	100.0%
50	random	30	26.7%	26.7%	20.0%	73.3%
50	hard	30	40.0%	23.3%	16.7%	60.0%
50	easy	30	23.3%	23.3%	16.7%	76.7%
50	topk_only	30	33.3%	30.0%	13.3%	66.7%
50	none	30	3.3%	-	-	96.7%
500	random	30	20.0%	13.3%	6.7%	80.0%
500	hard	30	33.3%	16.7%	6.7%	66.7%
500	easy	30	13.3%	13.3%	6.7%	86.7%
500	topk_only	30	30.0%	30.0%	13.3%	70.0%
500	none	30	16.7%	-	-	83.3%

(pick=non-refusal rate; in_topk=top-1 picked is one of the FAISS top-5; top1=ret=top-1 picked is the FAISS top-1; refusal=model emitted NONE.)

Headline findings

1. **Corpus-task mismatch is fundamental.** Even topk_only at sz=5 - a model handed exactly the FAISS-most-similar 5 ClawHub skills with no noise - refuses 67% of the time. ClawHub does not actually cover most TB task domains; the model knows it. 2. **`agree\_top1` <= 20% across all conditions.** When the model does pick, it rarely picks the FAISS top-1; it has independent judgement about which retrieved candidate (if any) is right. 3. **Hard noise increases pick rate at sz=50 (40%) over topk_only (33%). Hard-neg distractors look superficially more relevant than the actual FAISS top-K, so the model is more willing to commit. At sz=500 the same effect: hard 33% pick rate beats random 20% and easy 13%.* 4. *No-GT hallucination at sz=500: 16.7%.** Same pattern as SkillsBench but lower magnitude (SB was 18%) - TB's task descriptions are more specific, helping refusal even in pure-noise pools.

Implications

* The TB testset's "validated GT" is genuinely empty for most tasks; ~67% of TB tasks have no ClawHub skill the model would accept. This justifies P4's drop of incremental injection - it would have produced "no skill helps" labels for the same ~67% regardless of execution.

* For the controlled-study paper, TB serves best as a **negative control / OOD corpus**: the model's selection collapse on SkillsBench shows it can fail even when GT exists; on TB it shows the model correctly refuses when GT *doesn't* exist (mostly). Together that's a stronger story than either alone.

Caveats

- * 30 tasks (12.4% of TB) - full 241 should be run for paper-grade numbers.
- * Single seed.
- * "topk_only" pool ordering is rank-preserved (not shuffled); other modes shuffle. Need a topk_only_shuffled variant to verify ranking signal isn't just position bias.

2026-05-04: Terminal-Bench validated GT via Opus LLM-as-judge

Run: testsets/runs/2026-05-04-tb-judge.jsonl -> testsets/data/terminal_bench_validated.jsonl
Method: for each TB task x top-5 ClawHub FAISS retrieval = 1205 (task,skill) pairs, ask Opus 4.7 (Max-plan claude -p, conc=4): "would loading this skill materially help an autonomous shell agent solve the task?". YES -> ground-truth skill; tasks with all-NO are dropped.

Replaces the originally-spec'd "incremental skill injection via harbor" GT-construction loop. Harbor was ruled out because (a) harbor's terminus-2-skills agent calls models via LiteLLM/API only (no Max-plan support), so 241 x 6 task-runs would have cost ~\$700-\$1.7K; (b) per `feedback_retrieval_isolation.md`, agent-execution-as-GT was previously forbidden. LLM-as-judge keeps GT construction in pure forward-pass methodology while still labelling at the (task, skill) granularity.

metric	value
(task, skill) pairs judged	1205 (241 x 5)
YES verdicts	98 (8.1%)
NO verdicts	1106 (91.8%)
Unparseable (Anthropic AUP refusal)	1 (treated as NO)
Tasks kept (≥ 1 YES)	**62 / 241 (25.7%)**
Tasks dropped	179 (74.3%)
Multi-GT tasks (≥ 2 YES)	22
GT-count distribution	1: 40, 2: 12, 3: 7, 4: 2, 5: 1
YES distribution by FAISS rank	r1: 33, r2: 21, r3: 14, r4: 18, r5: 12
Wall-clock	~25 min (Opus 4.7 via Max plan)
Cost	\$0 marginal (Max plan)

Headline findings

1. **74% TB-task drop rate confirms the corpus-task mismatch.** ClawHub does not cover most terminal-bench task domains; even with strict top-5 retrieval, only 1 in 4 TB tasks has a single ClawHub skill that passes a strict-relevance bar. Consistent with the 67% refusal rate observed in the 2026-05-03 prompt-and-parse pilot - the same coverage gap surfaces under both methodologies. 2. **YES verdicts span all 5 ranks (33/21/14/18/12), not concentrated at rank 1.** FAISS embedding similarity has signal but is not a clean ranker - 65 of 98 GT skills are at rank ≥ 2 . This is exactly the dataset shape needed to evaluate whether the eval-target model's selection beats raw retrieval ordering. 3. **Multi-GT structure (22 tasks have 2-5 GT skills)** preserves the recall-vs-precision trade-off; the eval harness's `validated_gt_slugs` field can be scored as a set, not just top-1.

Caveats

- * Single judge model (Opus 4.7). For a paper-grade label, run a second-model agreement check (Sonnet or non-Anthropic) and treat the intersection as high-precision GT, the union as recall-leaning GT.
- * Strict prompt may over-reject niche but valid skills. Spot-checking 10 NO verdicts before scaling decisions.
- * Judge sees only skill name + ~200-char description from the corpus, not the full SKILL.md (the corpus's `data/raw/skills/_listing.jsonl` only stores summaries). Adding full-text retrieval would likely raise the YES rate.
- * 62-task validated TB testset is small. Combine with the 88-task SkillsBench validated set for a 150-task evaluation pool when running the prompt-and-parse harness end-to-end.

2026-05-04: Terminal-Bench validated-grid eval (point 4c/5)

Run: `testsets/runs/2026-05-04-tb-validated-grid.jsonl` (summary .summary.json) Config: 62 validated TB tasks x {5, 50, 500} pool size x {random, hard, easy, topk_only, none} noise x seed 0 = **930 trials, 0 errors**. Backend: `claude -p Max plan`, model `claude-sonnet-4-5`, conc 6 (initial) / 4 (after rate-limit resume), wall-clock ~80 min. Pool builder: GT slug pulled from `validated_gt_slugs`; multi-GT scoring (Hit@1 = top-1 in GT set; Recall = |picked ? GT| / |GT|; MRR over picked-list).

size	noise	n	Hit@1	Recall	MRR	refusal	pick	in_topk
5	random	62	74.2%	58.1%	0.742	12.9%	87.1%	87.1%
5	hard	62	69.4%	54.1%	0.694	17.7%	82.3%	82.3%
5	easy	62	75.8%	58.9%	0.758	17.7%	82.3%	82.3%

5	topk_only	62	74.2%	57.6%	0.742	16.1%	83.9%	83.9%
5	none	62	-	-	-	**100.0%**	0.0%	0.0%
50	random	62	75.8%	59.5%	0.758	12.9%	87.1%	85.5%
50	hard	62	54.8%	45.1%	0.548	19.4%	80.6%	59.7%
50	easy	62	75.8%	57.0%	0.758	17.7%	82.3%	82.3%
50	topk_only	62	**80.6%**	64.3%	0.806	11.3%	88.7%	88.7%
50	none	62	-	-	-	96.8%	3.2%	0.0%
500	random	62	58.1%	42.8%	0.581	30.6%	69.4%	62.9%
500	hard	62	**41.9%**	34.2%	0.419	27.4%	72.6%	43.5%
500	easy	62	48.4%	36.2%	0.484	40.3%	59.7%	56.5%
500	topk_only	62	79.0%	63.5%	0.798	17.7%	82.3%	82.3%
500	none	62	-	-	-	83.9%	16.1%	0.0%

Key takeaway

The model underperforms a trivial embedding-argmax baseline under hard distractors, despite the correct skill being among the top-5 candidates. This indicates a failure of fine-grained discrimination rather than retrieval. See figures/collapse_analysis.pdf for the embedding-neighborhood diagnostics (Figure 1).

baseline / picker	Hit@1
random-from-top-5 (no LLM, uniform)	31.3%
LLM @ sz=500 hard	**41.9%**
LLM @ sz=500 easy	48.4%
embedding-argmax oracle (no LLM, deterministic)	**53.2%**
LLM @ sz=50 topk_only (clean condition)	80.6%

Wrong picks concentrate near the task in embedding space: 24.7% are the corpus argmax itself, 71.9% are in the corpus top-10, median rank = 4 of 44,787. 49% of wrong picks satisfy $|\text{sim}(\text{GT}, \text{task}) - \text{sim}(\text{picked}, \text{task})| < 0.05$; 28% of wrong picks have $\text{sim}(\text{picked}) > \text{sim}(\text{GT})$.

Headline findings

1. **Validated GT recovers a real signal:** Hit@1 in the small-pool / topk_only setting (74-81%) is now in the same ballpark as SkillsBench's GT-aware grid (76-85% at sz=5 random) - confirming the Opus-judge GT is a meaningful target, not noise. 2. **Selection-collapse pattern reproduces on TB:** hard-noise Hit@1 falls 69.4% (sz=5) -> 54.8% (sz=50) -> **41.9%** (sz=500) - a 27.5 pp drop, the same shape as SkillsBench (59.6% -> 41.6%). The collapse-under-hard-noise phenomenon is now a *cross-dataset* finding. 3. **topk_only is the stability ceiling:** at sz=50 (80.6%) and sz=500 (79.0%) it nearly equals sz=5 (74.2%) - when the top-K ordering itself is preserved (no shuffled noise), model handles huge pools fine. Confirms the failure mode is *distractors mixed near GT*, not pool size per se. 4. **No-GT refusal degrades with pool size:** 100% (sz=5) -> 96.8% (sz=50) -> 83.9% (sz=500). Identical to SkillsBench (97.8% / 95.5% / 82.0%). The 16% sz=500 hallucination rate is robust across both datasets - a real retrieval-system safety concern. 5. **Recall mostly tracks Hit@1** because most validated tasks are single-GT (40 of 62); multi-GT tasks (22) are where Recall lags Hit@1 (e.g., sz=50 topk_only: 80.6% Hit@1 but 64.3% Recall).

Caveats

- * Single seed; need {0,1,2} replication for variance bars.
- * 62 tasks limits statistical power vs. SkillsBench's 89.
- * "topk_only" pool ordering is rank-preserved (not shuffled). Position bias not isolated; could be a follow-up.
- * Initial run hit Max-plan rate limit at 747/930 (5:40pm Indianapolis); resumed at conc=4 with --resume and finished cleanly.

Dataset construction note (read before citing any numbers)

Ground-truth skills are selected via an LLM judge (Opus 4.7) from the top-5 embedding-retrieved candidates per task. **All results are therefore conditioned on the correct skill being highly retrievable** - i.e., within the top-5 by embedding similarity to the task description. Concretely, of the 62 validated tasks, 33 have GT-rank-in-corpus = 1 and 29 have GT-rank-in-corpus in [2, 5]; none lie deeper than rank 5 by construction.

Findings should be interpreted as: **even under this favorable retrievability regime, the model fails to reliably select the correct skill under hard distractors**. The 179 tasks dropped at the GT-construction step are out of scope for the validated-grid claims; they are characterised separately in the 2026-05-03 TB-pilot findings as a coverage-gap / refusal-calibration story.